



Enrique Abril Contreras

AI Engineer · Multi-Agent Systems · RAG · LLM Apps in Production

📍 LATAM · Remote · ✉ kiketachira@gmail.com · 📞 +57 300 105 7454

🌐 [linkedin.com/in/enrique-abril-contreras](https://www.linkedin.com/in/enrique-abril-contreras) · 🐙 github.com/kiketach · 🌐 bio.fyllu.net

SUMMARY

AI Engineer with proven experience taking AI solutions **from design to real production**: conversational agents, multi-agent systems, and RAG architectures connected to business processes. **Former CTO** of a multi-tenant AI SaaS platform running in production with real clients. I build the full stack of an AI application: agent orchestration (Agency Swarm, OpenAI Agents SDK, Google ADK), a multi-provider LLM layer (LiteLLM), tools/MCP, guardrails and structured outputs, RAG over private data (pgvector), real-time voice (Gemini Live / Vertex AI), FastAPI backends, and cloud deployment (Cloud Run, Railway).

PROFESSIONAL EXPERIENCE

Full Stack AI Engineer & Technical Account Lead — MIMEIA

2025 – PRESENT

Conversational AI solutions for multiple clients on Symphonai, the in-house (DDD) platform for WhatsApp-first agent apps, with contributions to the core.

- Conversational agents and CRMs over WhatsApp for clients across verticals (pettech, commercial consulting, retail, aesthetics/health clinic) as *thin consumers* of the platform.
- **Subscription-tier agent** (pettech): instructions, tools, and guardrails composed at runtime based on the user's tier; input/output guardrails with a lightweight classifier (~100 ms) to block medical advice; multimodal UX (text, voice, image, PDF).
- **AI-native CRM** (commercial consulting): **event-sourced** design (immutable event ledger; pipeline stages derived by rules, not stored) with multi-agent orchestration and meeting scheduling.
- **Booking vertical** (*capacity-aware* appointment ERP): slots derived on-demand and cached in Redis with Postgres as the source of truth; an agent acting only through tools that wrap domain services, guaranteeing capacity invariants.
- **Platform core**: multi-provider AI orchestration (LiteLLM + OpenAI Responses API) with normalized token/cost accounting (OpenAI, Anthropic, Google, ElevenLabs); an event-driven backbone on **Redis Streams**; decisions documented in ADRs.
- **Resilient LLM routing**: Gemini on Vertex AI (service account) with automatic fallback to another provider on errors or rate limits.

CTO & Lead AI Engineer — RentOSO SpA

2025 – 2026

CRM + AI platform for the real estate sector, multi-tenant SaaS model. Chile · Remote.

- Led the architecture and development of the platform, as well as the development team.
- AI agents in production with **Google ADK + Vertex AI Agent Engine + Gemini**, on a **multi-tenant architecture with no redeployment** (onboarding a client = configuration, not a deploy).
- Multi-tenant on Supabase + PostgreSQL with **Row-Level Security**; domain separation: FastAPI (AI and integrations) and NestJS (transactional core).
- Conversational workflows with document validation, lead scoring, and WhatsApp lead-capture automation.
- Integrations: WhatsApp Business API (Meta Tech Provider), Mercado Libre Inmuebles, Twilio (voice), Floid (financial scoring), Cal.com.
- **Impact: 11+ agencies live in production**; negotiated partnerships with listing portals and with Mercado Libre Inmuebles to distribute the AI agents.

AI Engineer · Intelligent Document Systems — Audisoft

2024 – 2025

AI for document automation and internal audit, enterprise clients (Ecopetrol, Afinia, Banamex). Remote.

- Conversational assistants with **semantic retrieval (RAG)** over unstructured documents (PDF, email, images, Office).
- **ETL** pipeline from SQL Server into vector stores for semantic search; hybrid SQL + NoSQL systems for enterprise retrieval.
- Integration of OpenAI, Azure Document Intelligence, and MongoDB Atlas; decoupled architecture (processing, indexing, semantic query, generation).

PERSONAL PROJECT

KikeMind — AI platform for developer growth

IN PRODUCTION

End-to-end design, development, and deployment: product + architecture + code. [kikemind-9420c.web.app](#)

- **Conversational technical mentor** (text and voice) on system design, cloud/GCP,DDD, agents/LLMs, testing, and interviews; English suite (write, speak, listen, read).
- **Real-time voice with Gemini Live** (native audio) on Vertex AI, over a direct WebSocket to Cloud Run with a single-use ticket.
- Clean Architecture / SOLID with a port-adaptor-decoupled AI layer; **FastAPI · SQLModel · Supabase (JWT) · Alembic · React 19**; deployed on Firebase Hosting + Cloud Run.

CommerceEngine / Fyllu — CRM + AI agent + POS for WhatsApp conversational commerce

PILOT IN PRODUCTION

Own product for LATAM SMBs · pilot client Le Boutique · Cloud Run · [fyllu.net](#)

TEACHING

Guest Instructor — Universidad de Santiago de Chile (USACH), Faculty of Technology

AUG 2026

Seventh FACTER Extension School · Live online.

Designed and delivered the course “**Build Your First AI Agent: From Idea to a Working Assistant**” — 12 hours, 51 participants, mostly professionals with no prior programming experience. From zero to a working agent: **function calling** with custom tools, **MCP (Model Context Protocol)** connecting external servers, question answering over their own PDFs, and a web interface with voice.

51

participants

100%

rated command of
the subject very good

100%

learned something
of genuine value

84%

course fully met
their expectations

TECHNICAL SKILLS

AI / Agents: Multi-agent systems (MAS), Agency Swarm, OpenAI Agents SDK, Google ADK, PydanticAI, LiteLLM, MCP (Model Context Protocol), RAG, embeddings, vector databases, function/tool calling, structured outputs, guardrails, prompt engineering, evaluation

Models & voice: OpenAI (GPT), Anthropic (Claude), Google Gemini, Gemini Live (Vertex AI), ElevenLabs, open-source models; multimodal (text/image/voice)

Backend: Python, FastAPI, asyncio, Pydantic, SQLModel, TypeScript, NestJS, Next.js

Architecture: DDD, Clean Architecture / ports-adapters, multi-tenant, event-driven (Redis Streams), SSE/WebSockets, platform-contract (thin consumers)

Data / Vector: PostgreSQL, Supabase, pgvector, Row-Level Security, MongoDB Atlas, SQL Server, Redis

Cloud / DevOps: Google Cloud Platform (Cloud Run, Vertex AI, Secret Manager), Railway, Docker, CI/CD, Firebase, Azure

AI-assisted dev: Claude Code, Codex, Devin (SWE), DeepSeek Harness — agentic workflows, automated change review

Integrations: WhatsApp Business API (Wappa), Mercado Libre API, Twilio, Meta API/CAPI, Google Calendar, webhooks, OAuth

EDUCATION

- ✓ **U. Nac. Experimental del Táchira** — Computer Engineering · 2010–2012
- ✓ **Generation Colombia** — Full Stack Development (Java) · 2023
- ✓ **BIT Bogotá** — Front-end / Web Applications · 2022–2023

CERTIFICATIONS & LANGUAGES

- ✓ **Microsoft Azure AI Fundamentals (AI-900)**
- ✓ **Microsoft Azure Data Fundamentals (DP-900)**
- ✓ Spanish (native) · English (B2 — upper intermediate)